



BE PART OF WHAT'S NEXT

← ALL JOBS



Distributed Systems Engineer (L5) - Data Platform

USA - Remote

APPLY NOW

Job Posting Date

02-13-2026

Job Requisition ID

JR28718

Teams

Engineering

Work Type

Remote

At Netflix, our mission is to entertain the world. Together, we are writing the next episode - pushing the boundaries of storytelling, global fandom and

making the unimaginable a reality. We are a dream team obsessed with the uncomfortable excitement of discovering what happens when you merge creativity, intuition and cutting-edge technology. Come be a part of what's next.

The data platform teams at Netflix enable us to leverage data to bring joy to our members in many different ways. We provide centralized data platforms and tools for various business functions at Netflix, so they can utilize our data to make critical data-driven decisions. We do all the heavy lifting to make it easy for our business partners to work with data efficiently, securely, and responsibly. We aspire to lead the industry standard in building a world-class data infrastructure, as Netflix leads the way to be the most popular and pervasive destination for global internet entertainment.

We are looking for Distributed Systems Engineers to help evolve and innovate our infrastructure. We are committed to building a diverse and inclusive team to bring new perspectives as we solve the next set of challenges. In addition, we are open to remote candidates. We value what you can do, from anywhere in the U.S.

Spotlight on Data Platform Teams:

Data Discovery and Governance

Data Discovery and Governance are strategic bets at Netflix that underpin our ability to leverage data as a valuable asset. The Netflix data discovery and governance team owns the foundational building blocks necessary for ensuring all our data is used effectively, efficiently, securely, and in compliance with regulations. Specifically, this allows Netflix to organize its data better, take informed actions with increased precision in security, privacy, and efficiency, improve our risk management capabilities, increase confidence in the quality of the data, and enable the discovery & auditing of high-value datasets.

While Netflix has built a strong data platform foundation, we are still in the very early stages of setting up a scalable and extensible data discovery platform to help the company derive higher-level insights and values of data across new businesses like Games, Live, Ads, on top of our streaming service.

Specifically, this team owns the Netflix-wide data catalog to capture and infer business metadata across all datasets at Netflix, distributed search and lineage infrastructure to effectively discover data and track lineage across different

systems; and an extensible policy engine framework that allows our stakeholders to customize data policy rules for all datasets.

Online DataStores - KeyValue Team

The KeyValue team (part of the Online Data Stores portfolio) offers a managed key value store to meet Netflix's operational data requirements across all lines of business. Our focus is on developing and maintaining a high-performance, reliable, and efficient key-value store. Additionally, we enhance developer productivity by providing secure, intuitive and opinionated access layer for our backing storage engines including Cassandra.

This team is passionate about building innovative, efficient, performant, and scalable offerings for our customers. Our team also engages with the Cassandra open source community, actively contributing to and responsibly stewarding the development and use of open source software.

Online DataStores - Caching

Our focus is on developing and maintaining highly performant, reliable, and cost-efficient data storage solutions at Netflix's extreme scale. These data stores are built from a combination of world-class in-house solutions that are open-sourced and used by top technology firms, as well as state-of-the-art data solutions from public cloud offerings. Additionally, we enhance developer productivity by providing opinionated and intuitive, secure abstraction layers on top of the base storage systems.

The Online Data Stores team is a customer-obsessed infrastructure team that strives to accelerate the growth and diversification of Netflix's business verticals through world-class data storage technologies. Our team also seeks to impact the greater technology landscape by engaging with open source communities, actively contributing to, and responsibly stewarding the development and use of open source software.

The Netflix Online Data Stores Caching Team provides extremely low-latency yet high-throughput in-memory and distributed data storage solutions across the company. Our services power mission-critical Tier-0 Netflix services such as the streaming Catalog and Content Delivery Network (CDN) metadata. We are looking for passionate and innovative software engineers to join our team to build novel caching architectures that push the limits of performance, cost efficiency, utility, and usability.

Big Data Orchestration

The Big Data Orchestration team at Netflix empowers data engineers, data scientists, and ML researchers to create and manage computing workflows that leverage the full power of the Netflix Data Platform. This team is central to data processing at Netflix, hosting all ETL and batch ML data workloads running across domains as diverse as homepage personalization, studio & content application, and critical company reporting tasks. We provide both self-service and managed offerings in ETL and DAG workflow creation and execution, along with a variety of supporting services for failure robustness, data validation, and featureful user interfaces. We solve challenging problems working with data at scale, providing world-class uptime and efficiency features, and supporting Netflix engineers to delight our members, advance research frontiers and achieve ambitious business goals.

Analytics Infrastructure and Enablement

Our Analytics Infrastructure and Enablement team offers a platform that enables data practitioners to develop and manage data models for broad analytical use and to produce analytic products with efficiency and ease. The ultimate goal of the team is to enable broader self-service of high-fidelity analytics by developing robust APIs, user-friendly interfaces, and semi-automated data management and governance systems. Our team offers abstractions on Netflix's expansive data warehouse, and other data platform components, to provide easy access to semantic data models and easier development of analytical reporting and visualizations.

Data Security & Infrastructure

The Data Security and Infrastructure (DSI) team acts as a platform for our own data platforms, providing foundational infrastructure and tooling that empower Netflix to innovate rapidly in delivering cutting-edge data and analytics systems. Instead of building custom scaffolding for every new system, we focus on creating scalable, high-leverage infrastructure tailored to the demands of running data systems at Netflix's scale.

Our work spans infrastructure, security, deployment systems, and more, ensuring the reliability, security, and performance of our platform. A unique aspect of our team is our diverse tech stack, with many of our tools and systems developed in Go, Python and Java. If you thrive on solving complex problems, enjoy working with multiple programming languages, and want to make an impact at scale, this team could be a perfect fit for you!

Data Developer Experience (DDX)

The Data Developer Experience (DDX) team at Netflix is dedicated to enhancing the experience of data practitioners—including Data Engineers, Analytics Engineers, ML Engineers, and Data Scientists. As we evolve toward the next generation of our data platform, the DDX team is focused on boosting productivity by developing cutting-edge data tools and infrastructure that eliminate the friction and toil of building, maintaining, and leveraging data products. If you're passionate about shaping the future of developer experience, making it faster and easier for Netflix data practitioners to deliver outstanding data products, come join our team!

This would be your dream job if you enjoy:

- Solving real business needs at a large scale by applying your software engineering and analytical problem-solving skills.
- Architecting and building a robust, scalable, and highly available distributed infrastructure.
- Leading cross-functional initiatives and collaborating with engineers, product managers, and TPM across teams.
- Sharing our experiences with the open source communities and contributing to Netflix OSS.

About you:

- 7+ years experience in crafting complex, scalable distributed data infrastructure
- Proficiency in Java, C++, Golang, or Python with a solid understanding of multi-threading and memory management
- Proven track record of developing and maintaining high-impact systems
- Experience building and operating scalable, fault-tolerant, distributed systems
- You have a BS in Computer Science or a related field

Generally, our compensation structure consists solely of an annual salary; we do not have bonuses. You choose each year how much of your compensation you want in salary versus stock options. To determine your personal top of market compensation, we rely on market indicators and consider your specific job family, background, skills, and experience to determine your compensation